



# Mixed-Precision Attention Mechanisms for Efficient Transformers


## Fixed Window Attention



## Adaptive Mixed-Precision



## Block-Aligned Mixed-Precision



High Precision (FP16) Low Precision (FP8) Sparse (0-bit)